

Traversability Estimation for Ground Robots

Anonymous Author(s)

Affiliation

Address

email

1 **Abstract:** Mobile robots operating on unstructured terrain must predict which areas
2 of the environment they are able to pass in order to plan feasible paths and
3 to react to unforeseen terrain patterns. We address traversability estimation as a
4 heightmap classification problem: we build a convolutional neural network that,
5 given an image representing the heightmap of a patch of terrain, predicts whether
6 the robot will be able to traverse such patch from left to right. The classifier
7 is trained for a specific wheeled robot model (but may implement any other lo-
8 comotion type such as tracked, legged, snake-like), using simulation data on a
9 variety of procedurally generated training terrains; once trained, the classifier can
10 be quickly applied to patches extracted from unseen large heightmaps, in multiple
11 orientations, thus building oriented traversability maps. We quantitatively validate
12 the approach on real-elevation datasets and implement a path planning approach
13 that employs our traversability estimation.

14 **Keywords:** traversability estimation, convolutional neural networks

15 1 Introduction

16 In most indoor scenarios, the environment is trivially partitioned in traversable areas (clear floor)
17 and known obstacles (i.e. walls, objects, and other areas through which the robot can't pass); once
18 traversable and non-traversable areas are known, path planning (an essential task for mobile robots)
19 is solved using well-known algorithms [1].

20 In scenarios with uneven terrain, such as outdoors, segmenting the environment in traversable and
21 non-traversable areas is not as trivial. For example, some areas may be traversable only in a few
22 specific directions: a wheeled robot with limited power could be able to descend but not to ascend a
23 steep slope; a powerful, long and narrow robot with an high center of mass could be able to traverse
24 the same environment both uphill and downhill, but may capsize if traversing perpendicularly to
25 the steepest direction; a bicycle can hop up a side-walk as long as it is proceeding more or less
26 perpendicular to the step, but will crash when approaching the step from an oblique angle. Moreover,
27 criteria for traversability may not be intuitive or easy to model a-priori: a legged robot may be able
28 to negotiate very challenging terrain but get stuck on flat ground with deep holes of comparable size
29 as its feet; vacuum cleaner robots get stuck over power cords laying on the ground; a car with a low
30 clearance may get stuck on a speed bump.

31 We consider the problem of estimating where and in which directions a given 3D terrain is locally
32 traversable by a specific ground robot, using a general approach based on machine learning that
33 applies regardless on the robots locomotion method (wheeled, tracked, legged, snake-like), physi-
34 cal characteristics (size, motor torque), and low-level controller (anti-skid algorithms for wheeled
35 robots, foothold selection and gait selection algorithms for legged robots). In the following we de-
36 fine that a given portion of a terrain is *locally traversable* in a given direction if the robot, placed
37 in such position and direction, can proceed straight for at least a short distance when driven by its
38 low-level control algorithm.

39 When robot control parameters are known, traversability is only affected by the characteristics of
40 the terrain around the robot's position. Therefore, for a given robot pose X^{robot} with position p and
41 orientation θ , we consider as input an heightmap patch centered in p and rotated in such a way that
42 the robot is pointing towards the top of the patch. This information is represented as a heightmap

43 image where pixels indicate height values. Our task is cast as a binary image classification problem
44 (with classes *traversable* vs *non-traversable*), and solved by training a simple convolutional neural
45 network.

46 Training data is generated by simulating the robot on many procedurally generated training ter-
47 rains, which represent a variety of obstacles such as ramps, steps, bumps, holes and rugged areas;
48 during such simulations, the robot is spawned in random positions and orientations, and instructed
49 to proceed straight ahead; its progress is monitored and instances for traversable (where the robot
50 successfully proceeds) and non-traversable (where the robot can't proceed) heightmap patches are
51 continuously recorded. Once the model is learned from such training data, it can be applied densely
52 on any large unseen heightmap. Because accurate physical simulation is expensive, evaluating the
53 classifier on many points and orientations of a test terrain is orders of magnitude faster than simu-
54 lating the robot on these poses.

55 The **main contribution** of this work relies on the proposal of a complete framework for traversabil-
56 ity estimation detailed in Section 3. It includes i) a procedural generation of heightmaps, ii) a novel
57 approach for elevation-map based classification, iii) an orientation-based traversability representa-
58 tion, and iv) an application of such representation on path planning in real elevation maps and with
59 a real robot. Experimental validation and results are described in Section 4¹. Limitations and
60 extensions of our approach are discussed in Section 5.

61 2 Related Work

62 Estimating traversability is a fundamental capability for many animals and for autonomous mo-
63 bile robots, because most of their actions depend on their mobility. Traversability can be seen as
64 an *affordance* of the environment on the robot, or vice-versa [2], the former defines which terrain
65 characteristics allow the robot to traverse, the later includes robot's capabilities (e.g. morphology,
66 motor, sensing) to traverse the environment. In robotics, using a simplified (possibly learned) model
67 to estimate traversability is a common approach, because modeling the terrain, the robot and their
68 interaction accurately is difficult and expensive. Several approaches have been proposed to measure
69 traversability and to gather training data [3]. For instance, a robot may label a terrain as difficult
70 to traverse when sensing excessive vibrations [4, 5]; human experts can provide clues like pre-
71 ferred paths in a given terrain that avoid possible non-traversable regions [6]; the robot may learn a
72 traversability classifier on the spot, directly from experience [7].

73 The outdoor robotics community has researched traversability estimation from a variety of sensing
74 sources, such as stereo cameras [8], laser scanners [9], and time of flight cameras [10]. A common
75 approach to estimate traversability consists of two phases. First, from local sensory data, an elevation
76 map is derived [11]. Then, traversability estimates of parts of the elevation map are computed from
77 geometrical features like *slope*, *roughness*, and *step height* using pre-modeled functions specific to
78 the robot's locomotion and size [12].

79 Another major approach relies on the use of visual texture to classify the terrain type, e.g. dis-
80 criminating between rock, sand, and grass; then it derives a traversability score out of the assigned
81 label [4]. Terrain classification is a classical application for supervised machine learning techniques
82 based on generic visual features, transformation-invariant descriptors, texture features, or convo-
83 lutional neural networks (CNNs) [13, 14], which learn to extract problem-specific visual features.
84 DARPA's project *Learning Applied to Ground Robots project* [15] has advanced learning techniques
85 for terrain traversability classification, including some applications of deep learning [16].

86 3 Traversability Estimation Framework

87 Figure 1 illustrates the proposed framework for traversability estimation. Next, we describe how
88 the simulation environment is set-up (Section 3.1), detail the process for generating training and
89 evaluation datasets (Section 3.2), and finally describe classification approaches (Section 3.3).

¹Data and code to reproduce our results, video demonstrations, and additional information are anonymously available online: https://github.com/romarcg/traversability_estimation_pioneer

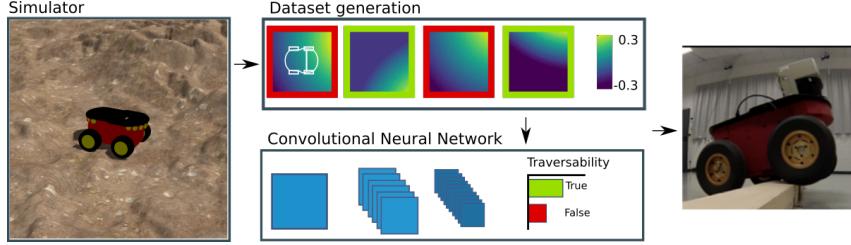


Figure 1: Robot model runs in simulation on procedurally generated terrains (left) to generate datasets linking heightmap patches with their traversability (top); on these datasets we train and evaluate classifiers to estimate the probability that a given heightmap patch is traversable or not (bottom). The learned classifier correctly predicts terrain traversability for a real robot (right).

90 3.1 Traversability from Simulation

91 We employ the Gazebo simulator and the ODE physics engine [17] for accurate physical simulation.
 92 We simulate the Pioneer 3-AT robot platform of size $49 \times 50 \times 29$ cm (illustrated in Figure 1)
 93 moving forward at constant velocity of $0.15m/s$ on an uneven terrain whose shape is determined by
 94 a heightmap. The robot’s trajectory on the terrain is captured to extract traversability informations.

95 3.1.1 Heightmap Generation

96 In order to generate meaningful training data, we need to simulate the robot moving on interesting
 97 terrains that pose varied and representative challenges. This could be achieved by using real
 98 heightmap data, by creating interesting terrains manually, and/or by synthesizing them using proce-
 99 dural generation techniques. We follow exclusively the third option. Each of our training terrains is
 100 generated as a grayscale heightmap, based on summing multiple realizations of random 2D simplex
 101 noise [18, 19]: a variant of Perlin noise [20], conceptually similar to it, frequently adopted in the
 102 procedural generation literature [21] with different periods.

103 For example, a heightmap, corresponding to simplex noise with period 30 cm and scaled such that
 104 it extends to a height of 20 cm, models a terrain with medium-sized smooth rocks while a period
 105 of 10 m with a 3 m height range yields a landscape with small steep hills. By summing the two
 106 heightmaps, one obtains rocky hills. If the *rocks* heightmap is multiplied by a value between 0 and
 107 1 which increases on the x coordinate, one gets flat hills for low values of x which becomes rocky
 108 hills for high values of x . Additional features such as holes, bumps protruding from a flat surface,
 109 rail-like indentations, steps, and others, can be generated by applying various functions on each
 110 value in the heightmap. Some example training terrains are represented in Figure 2. The size of
 111 each generated heightmap is 512×512 px that, when simulated, are scaled to represent surfaces
 112 of 10×10 m (≈ 2 cm/px resolution). In the supplementary material we provide source code for
 113 generating our training dataset and an appendix describing the approach in more detail.

114 3.1.2 Simulation Process

115 Once the simulation is initialized with a generated heightmap, the robot is set to a random pose on
 116 the map and moves forward at constant velocity without steering. After the robot reaches the edge of
 117 the map or gets stuck for some time, it is re-spawned to a different pose to generate a new trajectory.

118 For each trajectory, the robot poses and their corresponding heightmap patches are recorded; the
 119 heightmap patch associated to a pose is centred on the robot’s position and oriented in such a way
 120 that the robot is facing towards the right of the patch (see Fig. 1-right). If and only if the distance
 121 between the current pose $X^{\text{robot}}(t)$ and a future pose $X^{\text{robot}}(t + T)$ is greater than a threshold d and
 122 aligned with the robot’s orientation, then the patch is labelled as traversable.

123 Figure 3-left illustrates the traversability labelling for patches along a trajectory with $T = 1$ s and
 124 $d = 0.12$ m. Following this trajectory the robot is able to traverse a set of downward-sloped patches
 125 until it gets stopped by a bump almost as high as the robot itself. This behaviour is represented in
 126 Figure 3-right, where traversable patches are marked with green frames.

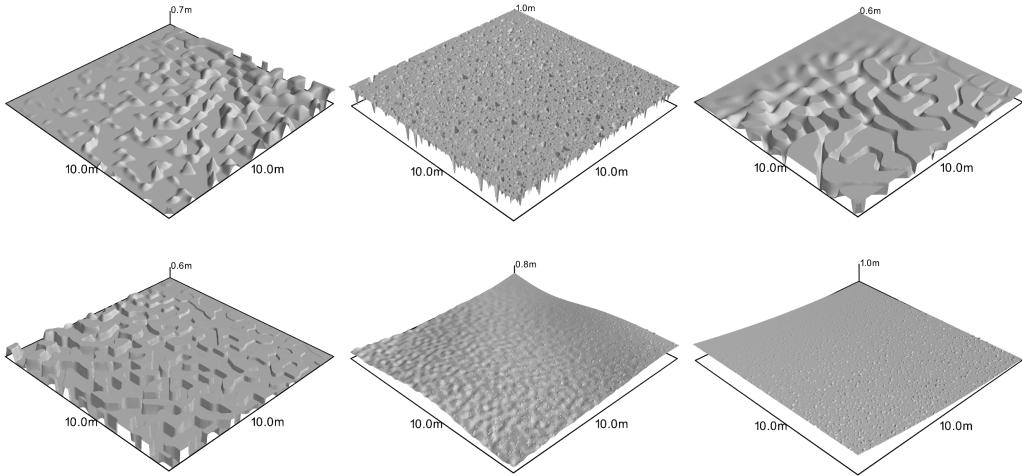


Figure 2: Examples of heightmaps for generating traversability datasets. These maps includes common terrain varieties as: bumps, holes, rails, steps, and slopes.

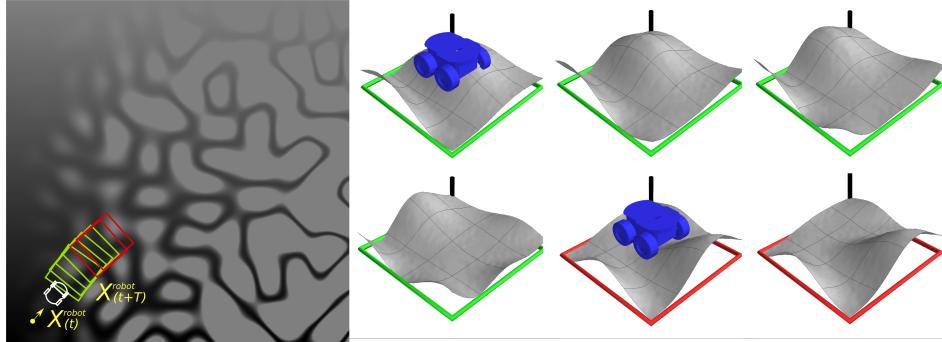


Figure 3: Trajectory extracted from a synthetic dataset (left). Robot silhouette and yellow arrow indicates the initial pose for the corresponding trajectory. Surfaces of the patches from the same trajectory (right). Green frames represent patches labelled as traversable, red frames as non-traversable.

127 3.2 Dataset Generation

128 From each simulated trajectory (on average 20 s), we sample the robot pose and the corresponding
 129 heightmap patch at 20 Hz. The patch and its traversability label represents an instance in the dataset.

130 We generate three datasets: D_{train} , a training dataset (450 k samples) from procedurally generated
 131 heightmaps (see Section 3.1.1); $D_{\text{eval,syn}}$, an evaluation dataset (150 k samples) with synthetic
 132 heightmaps also procedurally generated; $D_{\text{eval,real}}$, an evaluation dataset from real elevation maps
 133 (60 k samples). The maximum height of the heightmaps is 1 m for D_{train} and $D_{\text{eval,syn}}$; 3 m for
 134 $D_{\text{eval,gravelpit}}$. All datasets were generated using the robot described in Sec. 3.1, a patch size of
 135 60×60 px, $T \approx 1$ s and $d = 0.12$ m.

136 Figure 3-right illustrates patches from the trajectory of Fig. 3-left. The height values in each patch
 137 are offset in such a way that the patch center (i.e. at robot’s position) is mapped to height 0. This
 138 makes the patches independent on their absolute height on the heightmap, a feature that does not
 139 affect traversability.

140 Dataset $D_{\text{eval,gravelpit}}$ is generated using an elevation map from a Swiss gravel pit obtained by a flying
 141 drone [22]. The original area of this map is 0.48 km^2 , with a resolution of 5 cm/px and a maximum
 142 height of 50m. We cropped and scaled two regions of this map to form two heightmaps of 10×10 m
 143 with a resolution of 1.9 cm/px and a maximum height difference of 3 m. Figure 4 shows a reference
 144 image of the gravel pit and the surfaces of two extracted heightmaps.

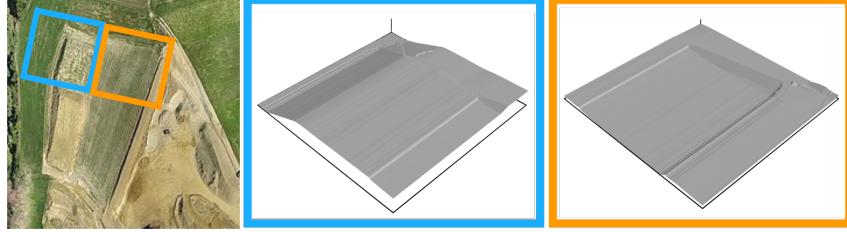


Figure 4: Surfaces of two heightmaps from real elevation data. Heightmaps were extracted from a mapping of a Swiss gravel pit [22]. Coloring represents height.

145 3.3 Training Traversability Classifiers

146 We address the problem of estimating terrain traversability as classification on heightmap data. We
 147 compare two alternative approaches: extracting descriptive features from each heightmap patch and
 148 then applying standard statistical classification techniques [23], or adopting Convolutional Neural
 149 Networks, a now-standard deep-learning approach [24] which operates directly on the raw input
 150 data. In either case, the output of the classifier indicates whether a patch is traversable.

151 For the **feature-based approach** we should extract from the input heightmap patch some quantities
 152 that provide traversability cues, for example the average terrain steepness in the robot’s motion
 153 direction (i.e. from left to right of the patch), or the maximum height of any steps in patch. In
 154 our case, we compute the Histogram of Gradients (HOG) of the heightmap patch, which includes
 155 these pieces of information (e.g., the gradient of a heightmap corresponds to the local steepness).
 156 Computing HOG over 6 orientations, 8×8 px per cell, and a block of 3×3 cells, results on a
 157 descriptor with 324 features that we classify by means of a Random Forest (RF) classifier [25] with
 158 10 trees.

159 In the **CNN-based approach**, it is expected that the network autonomously learns meaningful,
 160 problem-specific features; because the input data is high-dimensional and no prior knowledge of
 161 the problem is provided to the model, this approach requires more training data. Our CNN is built
 162 on the Keras [26] frontend powered by TensorFlow [27], and implements a 60×60 px input layer,
 163 followed by: a 3×3 convolution layer with 5 output maps; a 3×3 convolution layer with 5 output
 164 maps; a 2×2 Max-Pooling layer; a 3×3 convolution layer with 5 output maps; a fully connected
 165 layer with 128 output neurons; a fully connected layer with 2 output neurons followed by a softmax
 166 layer (output). All layers implement the ReLU activation function. The network is trained for 100
 167 epochs to minimize a categorical cross-entropy loss using the Adadelta optimizer.

168 4 Experimental Results

169 The following sections detail the comparison of the two classifiers described in Section 3.3 and an
 170 additional baseline dummy classifier that always returns the class most frequent in the training set.

171 4.1 Classification Results

172 Performance of the three estimators on the two different evaluation datasets (disjoint from the training
 173 set) $D_{\text{eval,syn}}$ and $D_{\text{eval,gravelpit}}$ is presented in Table 1. We observe that the CNN estimator out-
 174 performs both the baseline and feature-based approaches on both evaluation datasets. Performance
 175 is lower on the $D_{\text{eval,gravelpit}}$ dataset than on $D_{\text{eval,syn}}$, probably because of some elevation patterns,
 176 such as narrow and low barriers that may block the robot, which are not represented in procedurally-
 177 generated training data. Nonetheless, many features of real terrains such as slopes, rails and bumps
 178 are correctly classified.

179 4.2 Traversability Estimation on Real-world Datasets

180 We evaluate our trained classifiers on an additional real elevation map. This dataset consists on a
 181 mining quarry (see Fig. 5) of 0.51 km^2 [22]. The map contains some challenging roads and barriers

	$D_{\text{eval,syn}}$		$D_{\text{eval,gravelpit}}$	
	ACC	AUC	ACC	AUC
CNN	0.9264	0.9709	0.8327	0.8391
Feature-based	0.7031	0.7465	0.7698	0.8022
Baseline	0.5447	0.4971	0.5776	0.5044

Table 1: Performance of CNN, feature-based and baseline approaches. Accuracy (ACC) and area under the ROC curve (AUC) metrics are from a synthetic dataset $D_{\text{eval,syn}}$ generated similarly as the training dataset, and from the gravel pit dataset $D_{\text{eval,gravelpit}}$ described in Fig. 4.

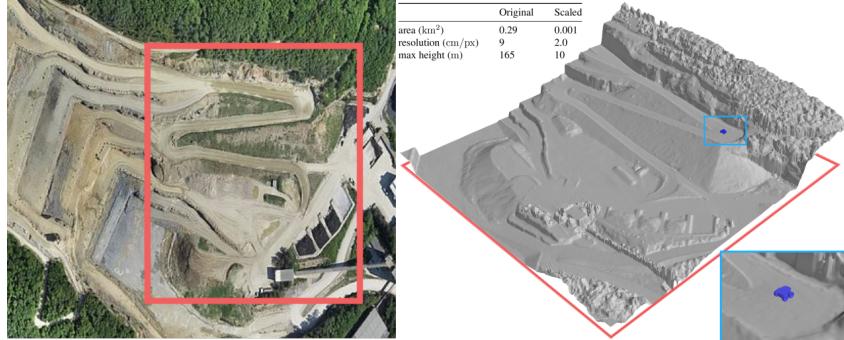


Figure 5: Heightmap extracted from Sensefly’s mining quarry dataset [22]. Reference top-view image and perspective view of the elevation map from the extracted red region are shown at left and right respectively. Pioneer 3-AT’s model is displayed for size comparison.

182 designed for cars and trucks. To make the map more suitable for our robot, which is about $0.1 \times$ the
183 size of a mining truck, we rescaled the map to $0.05 \times$ its original size.

184 For this analysis, we fix a direction and iterate over the entire heightmap extracting patches of
185 60×60 px with a stride of 5px. This process is equivalent as translating the robot’s position over the
186 map while keeping a fixed orientation. Figure 6 shows the traversability estimation for the mining
187 quarry for four orientations, indicated by the arrows. Traversability is represented as a colored
188 overlay on the surface of the heightmap (traversable is green, not traversable is gray). The minimum
189 traversability overlay (blue frame) is calculated as the the minimum traversability estimation at a
190 position from a set of possible orientations (in this case 32).

191 The estimator correctly marks the main road (≈ 1.9 m wide) as traversable in all directions. Very
192 steep or vertical slopes are recognized as non-traversable when going up but, dangerously, marked
193 traversable when going down. This particular result is congruent with the definition of traversability
194 we used, which only considers whether the robot can proceed but does not account for its safety.
195 Slopes are classified as traversable always downhill and, sometimes, transversally. Most of the
196 rough surface at the top of the elevation map is correctly found as non-traversable. The minimum
197 traversability representation highlights only areas that the robot can traverse in any of 32 orienta-
198 tions. These, correctly, correspond to roads and flat terrain.

199 4.3 Path Planning using Probabilistic Traversability Estimations

200 Traversability maps can be used to plan a path to a goal position which only passes through patches
201 which are traversable in the direction corresponding to the path’s local orientation. In the following,
202 we assume that the robot will get stuck (or crash, or end up in an unrecoverable state) if it attempts
203 to pass through a non-traversable patch².

204 We assume that traversability probabilities of non-overlapping patches are independent, i.e. that the
205 traversability along a path is Markovian. This allow us to compute the probability p to traverse a path

²this is a pessimistic assumption because this is not always the case: a robot attempting pass through a wall, for example, is unable to proceed but may be capable of turning back and taking another path

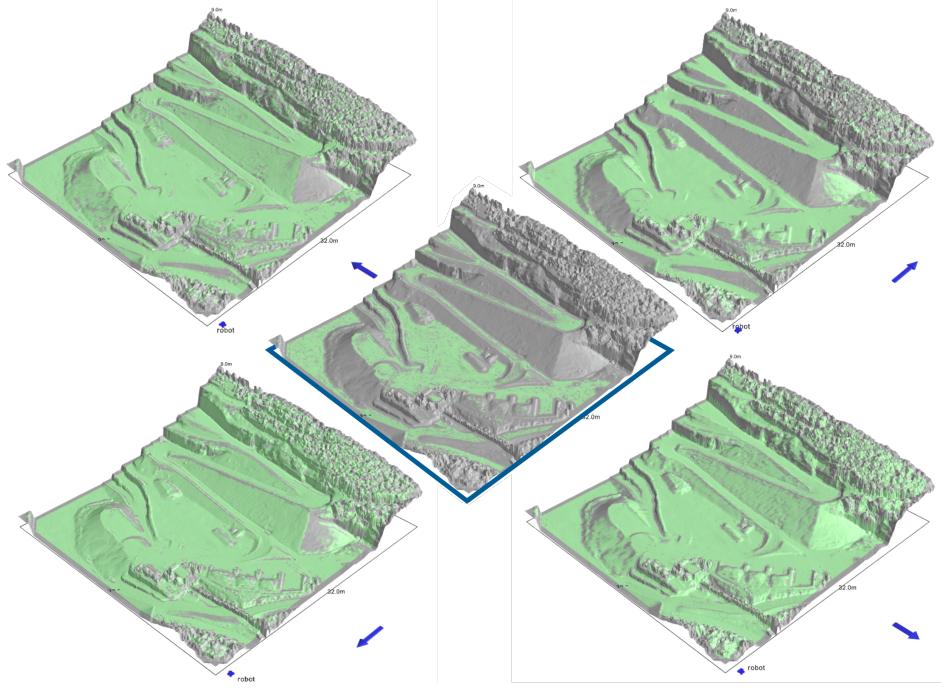


Figure 6: Traversability estimation for the mining quarry elevation map, for four orientations: 180° (left), 90° (up), 270° (down), 0° (right). Minimum traversability (blue frame) overlay over 32 possible orientations. For all figures, green indicates traversability and gray indicates non-traversability.

206 composed of segments (e_1, \dots, e_n) as $p((e_1, \dots, e_n)) = \prod_{i=1}^n p(e_i)$, where $p(e_i)$ is the probability
207 to traverse a segment.

208 We define a directed graph (N, E) where: nodes in N are placed on a regular square grid of side
209 length 18 cm, and; edges in E connect each node with its 8 nearest neighbors. For each edge $e \in E$
210 we compute length and traversal probability $p(e)$. $p(e)$ is the output of the traversability classifier
211 applied to a patch centered at the edge's origin and oriented along the edge direction (see Figure 7-
212 right).

213 Figure 7-left illustrates the solution of the multi-objective problem of computing the bests paths with
214 respect to length (to be minimized) and traversability (to be maximized). A rational agent would
215 choose among the set of Pareto-optimal paths. In this context, a path is Pareto-optimal if there exist
216 no alternative path that is both more traversable and shorter. The shortest path (represented in red
217 in the following) is Pareto-optimal, but will have in most cases a very low traversability probability.
218 The path with the maximum traversability (represented in green) will also be Pareto-optimal, but
219 may be unnecessarily long; between the two extremes, a potentially very large set of Pareto-optimal
220 paths exist, spanning the trade-off between traversability and length.

221 **Real robot experiment** We test how our approach applies to a real Pioneer 3-AT robot in an
222 environment where a pair of bars (1 m length \times 6 cm height) of width 6 cm and 8 cm respectively
223 lie horizontally aligned on a flat floor with a small gap of 15cm in-between (see Fig. 8); the bars
224 are fixed to the floor. We use an hand-generated heightmap representing the known geometry of the
225 scene. The robot is tasked to move from the area below the leftmost bar to the area above it. There
226 is no space to pass left of the bars, but plenty of space to pass on the right of the map. The robot is
227 constrained to move along vertical paths over (or beside) the obstacles. To verify whether a given
228 path is traversable, we teleoperate the robot following such path. We test three different trajectories:
229 the most direct one where the robot will always get stuck (red); the one when it may pass in the
230 space between the two bars by having wheels on one side climbing the bars, whereas the wheels
231 on the other side passing in the gap, this path has a probability of about 50% that was empirically
232 verified through several repetitions; and the longer but safest one going around the obstacles (green).

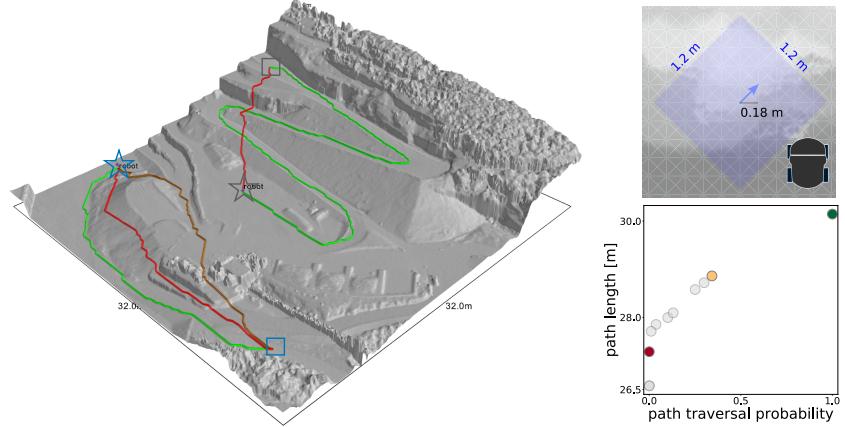


Figure 7: Left: a selection of Pareto optimal paths on the quarry map for two pairs of source (star) and target (square) locations. The paths are colored by traversal probability from red (non-traversable) to green (surely traversable). Right top: a portion of the planning graph on the quarry map. Nodes are placed on a regular grid of 18 cm. The blue edge’s traversal probability is estimated by applying the classifier on a 1.2 m x 1.2 m patch (light blue) centered at the edge origin and directed along the edge. Robot’s silhouette is shown for size comparison. Right bottom: The trade-off between path length and traversability for Pareto optimal paths between blue star and blue square. Colored dots correspond to paths drawn on the left figure.

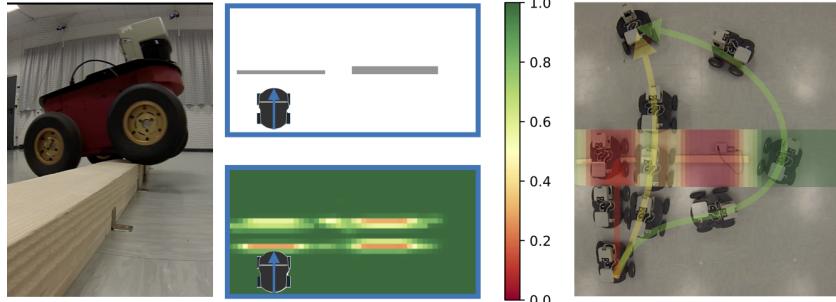


Figure 8: Left: Pioneer 3AT tries to pass over the bars. Center: height-map with two bars of width 6 cm and 8 cm (top) and patch traversability estimation moving upwards (bottom). Right: robot tries alternative routes. For all paths going straight upwards, the estimated traversal probability is displayed on top of the bars.

233 5 Conclusions and Perspectives

234 We presented a complete framework for traversability estimation that casts the problem as an
235 elevation-map classification task. Classifiers trained on simulation data capture complex charac-
236 teristics of different robot models and quickly estimate traversability maps on large unseen terrains.

237 We assumed that the terrain’s 3D shape is the only factor influencing traversability; other factors
238 may also be considered, such as compactness, friction, and instability. Our framework could handle
239 these factors provided they can be simulated: for example, if terrains whose 3D shape suggests a
240 loose gravel surface were simulated with lower friction, the classifier would automatically learn that
241 steep slopes can’t be negotiated on such surfaces; in this perspective, one may add additional inputs
242 to the classifier (such as the visual texture) which may help differentiate the actual surface.

243 Our current approach does not capture the robot dynamics, such as the speed with which the robot
244 approaches an obstacle: we limit our attention to slow robots operating on rugged terrains, where
245 dynamic aspects have negligible impact. We consider the addition of such information as we move
246 towards scenarios where complex locomotion is needed.

247 **References**

- 248 [1] R. Hadsell, M. Scoffier, U. Muller, Y. LeCun, and P. Sermanet. Mapping and planning un-
249 der uncertainty in mobile robots with long-range perception. In *In Proc. of the Int. Conf. on*
250 *Intelligent Robots and Systems (IROS)*, pages 1–6, 2008.
- 251 [2] U. E. and S. E. Traversability: A case study for learning and perceiving affordances in robots.
252 *Adaptive Behavior*, 18:258–284, 2010.
- 253 [3] V. Molino, R. Madhavan, E. Messina, A. Downs, S. Balakirsky, and A. Jacoff. Traversability
254 metrics for rough terrain applied to repeatable test methods. In *2007 IEEE/RSJ International*
255 *Conference on Intelligent Robots and Systems*, pages 1787–1794, Oct 2007. doi:10.1109/
256 IROS.2007.4399438.
- 257 [4] C. A. Brooks and K. D. Iagnemma. Self-supervised classification for planetary rover terrain
258 sensing. In *2007 IEEE Aerospace Conference*, pages 1–9, March 2007. doi:10.1109/AERO.
259 2007.352693.
- 260 [5] M. A. Bekhti, Y. Kobayashi, and K. Matsumura. Terrain traversability analysis using multi-
261 sensor data correlation by a mobile robot. In *2014 IEEE/SICE International Symposium on*
262 *System Integration*, pages 615–620, Dec 2014. doi:10.1109/SII.2014.7028109.
- 263 [6] D. Silver, J. A. Bagnell, and A. Stentz. Learning from demonstration for autonomous nav-
264 igation in complex unstructured terrain. *Int. J. Rob. Res.*, 29(12):1565–1592, Oct. 2010.
265 ISSN 0278-3649. doi:10.1177/0278364910369715. URL <http://dx.doi.org/10.1177/0278364910369715>.
- 266 [7] R. Hadsell, P. Sermanet, J. Ben, A. Erkan, J. Han, U. Muller, and Y. LeCun. Online learning
267 for offroad robots: Spatial label propagation to learn long-range traversability. In *Proceedings*
268 *of Robotics: Science and Systems*, Atlanta, GA, USA, June 2007. doi:10.15607/RSS.2007.III.
269 003.
- 270 [8] A. Chilian and H. Hirschmuller. *Stereo camera based navigation of mobile robots on rough*
271 *terrain*, pages 4571–4576. 2009.
- 272 [9] M. Herbert, C. Caillas, E. Krotkov, I. S. Kweon, and T. Kanade. Terrain mapping for a roving
273 planetary explorer. In *Proceedings, 1989 International Conference on Robotics and Automa-*
274 *tion*, pages 997–1002 vol.2, May 1989. doi:10.1109/ROBOT.1989.100111.
- 275 [10] H. Balta, G. D. Cubber, D. Doroftei, Y. Baudoin, and H. Sahli. Terrain traversability analysis
276 for off-road robots using time-of-flight 3d sensing. In *In Proc. of the Int. Workshop on Robotics*
277 *for Risky Environment - Extreme Robotics*, 2013.
- 278 [11] S. Lacroix, A. Mallet, D. Bonnafous, G. Bauzil, S. Fleury, M. Herrb, and R. Chatila. Au-
279 tonomous rover navigation on unknown terrains functions and integration. In *Experimental*
280 *Robotics VII*, ISER '00, pages 501–510, London, UK, UK, 2001. Springer-Verlag. ISBN 3-
281 540-42104-1. URL <http://dl.acm.org/citation.cfm?id=645627.757994>.
- 282 [12] M. Schwarz and S. Behnke. Local navigation in rough terrain using omnidirectional height. In
283 *ISR/Robotik 2014; 41st International Symposium on Robotics*, pages 1–6, June 2014.
- 284 [13] R. Hudjakov and M. Tamre. Aerial imagery terrain classification for long-range autonomous
285 navigation. In *2009 International Symposium on Optomechatronic Technologies*, pages 88–91,
286 Sept 2009. doi:10.1109/ISOT.2009.5326104.
- 287 [14] M. Längkvist, A. Kiselev, M. Alirezaie, and A. Loutfi. Classification and Segmentation of
288 Satellite Orthoimagery Using Convolutional Neural Networks. *Remote Sensing*, 8(4), 2016.
- 289 [15] L. D. Jackel, E. Krotkov, M. Perschbacher, J. Pippine, and C. Sullivan. The darpa lagr program:
290 Goals, challenges, methodology, and phase i results. *Journal of Field Robotics*, 23(11-12):945–
291 973, 2006. ISSN 1556-4967. doi:10.1002/rob.20161. URL <http://dx.doi.org/10.1002/rob.20161>.

- 294 [16] M. Shneier, T. Chang, T. Hong, W. Shackleford, R. Bostelman, and J. S. Albus. Learning
 295 traversability models for autonomous mobile vehicles. *Autonomous Robots*, 24(1):69–86,
 296 2008. ISSN 1573-7527. doi:10.1007/s10514-007-9063-6. URL <http://dx.doi.org/10.1007/s10514-007-9063-6>.
- 298 [17] N. Koenig and A. Howard. Design and use paradigms for gazebo, an open-source multi-robot
 299 simulator. In *2004 IEEE/RSJ International Conference on Intelligent Robots and Systems
 300 (IROS) (IEEE Cat. No.04CH37566)*, volume 3, pages 2149–2154 vol.3, Sept 2004.
- 301 [18] K. Perlin. Improving noise. In *ACM Transactions on Graphics (TOG)*, volume 21, pages
 302 681–682. ACM, 2002.
- 303 [19] M. Olano. Modified noise for evaluation on graphics hardware. In *Proceedings of the ACM
 304 SIGGRAPH/EUROGRAPHICS conference on Graphics hardware*, pages 105–110. ACM,
 305 2005.
- 306 [20] A. Lagae, S. Lefebvre, R. Cook, T. DeRose, G. Drettakis, D. S. Ebert, J. P. Lewis, K. Perlin,
 307 and M. Zwicker. A survey of procedural noise functions. In *Computer Graphics Forum*,
 308 volume 29, pages 2579–2600. Wiley Online Library, 2010.
- 309 [21] R. M. Smelik, K. J. De Kraker, T. Tutenel, R. Bidarra, and S. A. Groenewegen. A survey
 310 of procedural methods for terrain modelling. In *Proceedings of the CASA Workshop on 3D
 311 Advanced Media In Gaming And Simulation (3AMIGAS)*, pages 25–34, 2009.
- 312 [22] SenseFly. Elevation datasets. <https://www.sensefly.com/drones/example-datasets.html>, 2016. URL <https://www.sensefly.com/drones/example-datasets.html>.
- 314 [23] C. M. Bishop. *Pattern Recognition and Machine Learning (Information Science and Statistics)*.
 315 Springer-Verlag New York, Inc., Secaucus, NJ, USA, 2006. ISBN 0387310738.
- 316 [24] J. Schmidhuber. Deep learning in neural networks: An overview. *Neural Networks*, 61:85–
 317 117, 2015. ISSN 0893-6080. doi:<https://doi.org/10.1016/j.neunet.2014.09.003>. URL <http://www.sciencedirect.com/science/article/pii/S0893608014002135>.
- 319 [25] L. Breiman. Random Forests. *Machine Learning*, 45(1):5–32, 2001. ISSN 1573-0565. doi:
 320 [10.1023/A:1010933404324](https://doi.org/10.1023/A:1010933404324). URL <http://dx.doi.org/10.1023/A:1010933404324>.
- 321 [26] F. Chollet et al. Keras. <https://github.com/fchollet/keras>, 2015.
- 322 [27] M. Abadi, A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro, G. S. Corrado, A. Davis,
 323 J. Dean, M. Devin, S. Ghemawat, I. Goodfellow, A. Harp, G. Irving, M. Isard, Y. Jia, R. Joze-
 324 fowicz, L. Kaiser, M. Kudlur, J. Levenberg, D. Mané, R. Monga, S. Moore, D. Murray,
 325 C. Olah, M. Schuster, J. Shlens, B. Steiner, I. Sutskever, K. Talwar, P. Tucker, V. Van-
 326 houcke, V. Vasudevan, F. Viégas, O. Vinyals, P. Warden, M. Wattenberg, M. Wicke, Y. Yu,
 327 and X. Zheng. TensorFlow: Large-scale machine learning on heterogeneous systems, 2015.
 328 URL <http://tensorflow.org/>. Software available from tensorflow.org.